

# Database Implementation

## Database Implementation on GCP

### Connecting to GCP:

```
gcloud sql connect sp24-db-team065 --user=root;  
show databases;  
use AccommoSeek;  
show tables;
```

```
Welcome to Cloud Shell! Type "help" to get started.  
Your Cloud Platform project in this session is set to sp24-cs411-team065-dbmaster.  
Use "gcloud config set project [PROJECT_ID]" to change to a different project.  
chevady19980224@cloudshell:~ (sp24-cs411-team065-dbmaster)$ gcloud sql connect sp24-db-team065 --user=root;  
Allowlisting your IP for incoming connection for 5 minutes...done.  
Connecting to database with SQL user [root].Enter password:  
Welcome to the MySQL monitor.  Commands end with ; or \g.  
Your MySQL connection id is 358263  
Server version: 8.0.31-google (Google)  
  
Copyright (c) 2000, 2024, Oracle and/or its affiliates.  
  
Oracle is a registered trademark of Oracle Corporation and/or its  
affiliates. Other names may be trademarks of their respective  
owners.  
  
Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.  
  
mysql> show databases;  
+-----+  
| Database |  
+-----+  
| AccommoSeek |  
| TestModel |  
| information_schema |  
| mysql |  
| performance_schema |  
| sys |  
| test |  
+-----+  
7 rows in set (0.00 sec)  
  
mysql> use AccommoSeek;  
Reading table information for completion of table and column names  
You can turn off this feature to get a quicker startup with -A  
  
Database changed  
mysql> 
```

```
mysql> use AccommoSeek;  
Reading table information for completion of table and column names  
You can turn off this feature to get a quicker startup with -A  
  
Database changed  
mysql> show tables;  
+-----+  
| Tables_in_AccommoSeek |  
+-----+  
| City |  
| Hotel |  
| Precipitation |  
| Review |  
| Temperature |  
| crime_statistics |  
| hotel_review |  
| living_wage |  
+-----+  
8 rows in set (0.01 sec)
```

crime\_statistics, hotel\_review, living\_wage tables are used as temporary tables to clean and export data into City and Review table.

# Data Definition Language:

```
CREATE TABLE City (  
  Name VARCHAR(255) PRIMARY KEY,  
  State VARCHAR(255),  
  LivingWage REAL,  
  Population INT,  
  CrimeFrequency REAL  
);
```

```
CREATE TABLE Hotel (  
  Name VARCHAR(255),  
  Address VARCHAR(255),  
  CityName VARCHAR(255),  
  PRIMARY KEY (Name, CityName),  
  FOREIGN KEY (CityName) REFERENCES City(Name)  
);
```

```
CREATE TABLE Review (  
  UserName VARCHAR(255),  
  Title VARCHAR(255),  
  Text VARCHAR(1024),  
  Rating REAL,  
  Date datetime,  
  HotelName VARCHAR(255),  
  CityName VARCHAR(255),  
  PRIMARY KEY (UserName, HotelName, CityName),  
  FOREIGN KEY (HotelName) REFERENCES Hotel(Name),  
  FOREIGN KEY (CityName) REFERENCES City(Name)  
);
```

```
CREATE TABLE Temperature (  
  Month INT,  
  Year INT,  
  Temperature REAL,  
  CityName VARCHAR(255),  
  PRIMARY KEY (Month, Year, CityName),  
  FOREIGN KEY (CityName) REFERENCES City(Name)  
);
```

```
CREATE TABLE Precipitation(  
  CityState VARCHAR(255),  
  Month INT,  
  Precipitation REAL,  
  PRIMARY KEY (CityState, Month),  
  FOREIGN KEY (CityState) REFERENCES City(State)  
);
```

## Inserting Data

```
mysql> SELECT count(*) FROM Hotel;
+-----+
| count(*) |
+-----+
|      1854 |
+-----+
1 row in set (0.01 sec)

mysql> SELECT count(*) FROM Review;
+-----+
| count(*) |
+-----+
|      8492 |
+-----+
1 row in set (0.45 sec)

mysql> SELECT count(*) FROM Temperature;
+-----+
| count(*) |
+-----+
|       3100 |
+-----+
1 row in set (0.00 sec)
```

## Advanced Queries

### Advanced Query 1

**Advanced features:** join multiple relations, aggregation via GROUP BY, subqueries  
SELECT

c.Name AS CityName,  
ROUND(AVG(t.Temperature), 2) AS AverageTemperature,  
ROUND(AVG(c.LivingWage), 2) AS AverageLivingWage,  
ROUND(AVG(c.CrimeRate), 2) AS AverageCrimeRate,  
ROUND(AVG(r.Rating), 2) AS AverageRating

FROM City c

JOIN Temperature t ON c.Name = t.CityName  
JOIN Hotel h ON c.Name = h.CityName  
JOIN Review r ON h.Name = r.HotelName

GROUP BY c.Name

HAVING

AVG(t.Temperature) > (SELECT AVG(Temperature) FROM Temperature)  
AND AVG(r.Rating) > (SELECT AVG(Rating) FROM Review)

ORDER BY AverageRating DESC;

**Top 15 results:**

```
mysql>
mysql> (SELECT c.Name AS CityName,
-> ROUND(AVG(t.Temperature), 2) AS AverageTemperature,
-> ROUND(AVG(c.LivingWage), 2) AS AverageLivingWage,
-> ROUND(AVG(c.CrimeRate), 2) AS AverageCrimeRate,
-> ROUND(AVG(r.Rating), 2) AS AverageRating
-> FROM City c
-> JOIN Temperature t ON c.Name = t.CityName
-> JOIN Hotel h ON c.Name = h.CityName
-> JOIN Review r ON h.Name = r.HotelName
-> GROUP BY c.Name
-> HAVING AVG(t.Temperature) > (SELECT AVG(Temperature) FROM Temperature)
-> AND AVG(r.Rating) > (SELECT AVG(Rating) FROM Review)
-> ORDER BY AverageRating DESC);
+-----+
| CityName | AverageTemperature | AverageLivingWage | AverageCrimeRate | AverageRating |
+-----+
| Dallas | 19.79 | 25.46 | 6865.06 | 4.6 |
| Honolulu | 25.07 | 34.48 | 4248.94 | 4.36 |
| San Francisco | 15.33 | 35.55 | 9460.95 | 4.32 |
| San Antonio | 22.01 | 25.46 | 8206.13 | 4.26 |
| Sacramento | 17.55 | 35.56 | 6127.42 | 4.25 |
| Las Vegas | 17.66 | 26.15 | 6114.89 | 4.18 |
+-----+
6 rows in set (0.97 sec)
```

The output of the result has only 6 rows

## Command:

```
mysql> Explain Analyze (SELECT c.Name AS CityName,
->      ROUND(AVG(t.Temperature), 2) AS AverageTemperature,
->      ROUND(AVG(c.LivingWage), 2) AS AverageLivingWage,
->      ROUND(AVG(c.CrimeRate), 2) AS AverageCrimeRate,
->      ROUND(AVG(r.Rating), 2) AS AverageRating
-> FROM City c
-> JOIN Temperature t ON c.Name = t.CityName
-> JOIN Hotel h ON c.Name = h.CityName
-> JOIN Review r ON h.Name = r.HotelName
-> GROUP BY c.Name
-> HAVING AVG(t.Temperature) > (SELECT AVG(Temperature) FROM Temperature)
->      AND AVG(r.Rating) > (SELECT AVG(Rating) FROM Review)
-> ORDER BY AverageRating DESC);
```

## Analysis before indexing:

```
-----+
| -> Sort: AverageRating DESC (actual time=1482.110..1482.111 rows=6 loops=1)
|   -> Filter: ((avg(t.Temperature) > (select #2)) and (avg(r.Rating) > (select #3))) (actual time=1482.059..1482.074 rows=6 loops=1)
|     -> Table scan on <temporary> (actual time=1473.971..1473.991 rows=28 loops=1)
|       -> Aggregate using temporary table (actual time=1473.967..1473.967 rows=28 loops=1)
|         -> Nested loop inner join (cost=15151.34 rows=28624) (actual time=0.263..1179.819 rows=229700 loops=1)
|           -> Nested loop inner join (cost=5132.90 rows=5624) (actual time=0.147..72.835 rows=21200 loops=1)
|             -> Nested loop inner join (cost=1399.25 rows=3100) (actual time=0.130..31.011 rows=3100 loops=1)
|               -> Table scan on t (cost=314.25 rows=3100) (actual time=0.094..3.932 rows=3100 loops=1)
|                 -> Single-row index lookup on c using PRIMARY (Name=t.CityName) (cost=0.25 rows=1) (actual time=0.008..0.008
rows=1 loops=3100)
|               -> Covering index lookup on h using idx_hotel_cityname (CityName=t.CityName) (cost=1.02 rows=2) (actual time=0.0
07..0.013 rows=7 loops=3100)
|             -> Index lookup on r using HotelName (HotelName=h.'Name') (cost=1.27 rows=5) (actual time=0.020..0.051 rows=11 loops
=21200)
|           -> Select #2 (subquery in condition; run only once)
|             -> Aggregate: avg(Temperature.Temperature) (cost=624.25 rows=1) (actual time=1.758..1.759 rows=1 loops=1)
|             -> Table scan on Temperature (cost=314.25 rows=3100) (actual time=0.052..1.442 rows=3100 loops=1)
|           -> Select #3 (subquery in condition; run only once)
|             -> Aggregate: avg(Review.Rating) (cost=1802.75 rows=1) (actual time=6.249..6.249 rows=1 loops=1)
|             -> Table scan on Review (cost=953.75 rows=8490) (actual time=0.036..5.416 rows=8492 loops=1)
|
+-----+
```

## Attempt 1 : Index on Temperature

```
-----+
| -> Sort: AverageRating DESC (actual time=2269.986..2269.987 rows=6 loops=1)
|   -> Filter: ((avg(t.Temperature) > (select #2)) and (avg(r.Rating) > (select #3))) (actual time=2269.947..2269.959 rows=6 loops=1)
|     -> Table scan on <temporary> (actual time=2265.704..2265.721 rows=28 loops=1)
|       -> Aggregate using temporary table (actual time=2265.697..2265.697 rows=28 loops=1)
|         -> Nested loop inner join (cost=15151.34 rows=28624) (actual time=0.290..1117.020 rows=229700 loops=1)
|           -> Nested loop inner join (cost=5132.90 rows=5624) (actual time=0.201..68.954 rows=21200 loops=1)
|             -> Nested loop inner join (cost=1399.25 rows=3100) (actual time=0.181..30.232 rows=3100 loops=1)
|               -> Covering index scan on t using idx_temperature (cost=314.25 rows=3100) (actual time=0.123..3.926 rows=310
0 loops=1)
|                 -> Single-row index lookup on c using PRIMARY (Name=t.CityName) (cost=0.25 rows=1) (actual time=0.008..0.008
rows=1 loops=3100)
|               -> Covering index lookup on h using idx_hotel_cityname (CityName=t.CityName) (cost=1.02 rows=2) (actual time=0.0
06..0.012 rows=7 loops=3100)
|             -> Index lookup on r using HotelName (HotelName=h.'Name') (cost=1.27 rows=5) (actual time=0.018..0.048 rows=11 loops
=21200)
|           -> Select #2 (subquery in condition; run only once)
|             -> Aggregate: avg(Temperature.Temperature) (cost=624.25 rows=1) (actual time=0.846..0.846 rows=1 loops=1)
|             -> Covering index scan on Temperature using idx_temperature (cost=314.25 rows=3100) (actual time=0.036..0.640 rows=3100
loops=1)
|           -> Select #3 (subquery in condition; run only once)
|             -> Aggregate: avg(Review.Rating) (cost=1802.75 rows=1) (actual time=3.348..3.349 rows=1 loops=1)
|             -> Table scan on Review (cost=953.75 rows=8490) (actual time=0.023..2.750 rows=8492 loops=1)
|
+-----+
```

We expected to improve performance by indexing attributes in the HAVING clause, like 'Temperature'. The rationale behind this approach is that indexing should narrow down the search space to only include records meeting the specified condition, thereby improving efficiency. However, we've not seen a reduction in query costs. Specifically, even as the scan on 'Temperature' shifts from a table scan to an index scan, the cost remains static at 314.25. We suspect the reason is that our table's size is insufficient to exhibit significant improvements when weighed against the overhead of indexing in RDBMS.

## Attempt 2 : Index on CrimeRate

```
-----+
| -> Sort: AverageRating DESC (actual time=921.621..921.622 rows=6 loops=1)
|   -> Filter: ((avg(t.Temperature) > (select #2)) and (avg(r.Rating) > (select #3))) (actual time=921.585..921.596 rows=6 loops=1)
|     -> Table scan on <temporary> (actual time=917.378..917.390 rows=28 loops=1)
|       -> Aggregate using temporary table (actual time=917.375..917.375 rows=28 loops=1)
|         -> Nested loop inner join (cost=15151.34 rows=28624) (actual time=0.142..735.310 rows=229700 loops=1)
|           -> Nested loop inner join (cost=5132.90 rows=5624) (actual time=0.082..45.208 rows=21200 loops=1)
|             -> Nested loop inner join (cost=1399.25 rows=3100) (actual time=0.070..19.551 rows=3100 loops=1)
|               -> Table scan on t (cost=314.25 rows=3100) (actual time=0.049..2.636 rows=3100 loops=1)
|                 -> Single-row index lookup on c using PRIMARY (Name=t.CityName) (cost=0.25 rows=1) (actual time=0.005..0.005
rows=1 loops=3100)
|               -> Covering index lookup on h using idx_hotel_cityname (CityName=t.CityName) (cost=1.02 rows=2) (actual time=0.0
04..0.008 rows=7 loops=3100)
|             -> Index lookup on r using HotelName (HotelName=h.'Name') (cost=1.27 rows=5) (actual time=0.013..0.032 rows=11 loops
=21200)
|           -> Select #2 (subquery in condition; run only once)
|             -> Aggregate: avg(Temperature.Temperature) (cost=624.25 rows=1) (actual time=0.909..0.909 rows=1 loops=1)
|               -> Table scan on Temperature (cost=314.25 rows=3100) (actual time=0.036..0.721 rows=3100 loops=1)
|           -> Select #3 (subquery in condition; run only once)
|             -> Aggregate: avg(Review.Rating) (cost=1802.75 rows=1) (actual time=3.245..3.246 rows=1 loops=1)
|               -> Table scan on Review (cost=953.75 rows=8490) (actual time=0.024..2.719 rows=8492 loops=1)
|
+-----+
```

The addition of an index on the attribute present in the SELECT clause, specifically 'CrimeRate,' was not expected to influence the cost. The outcomes of our analysis align with these expectations, confirming that the index has no impact on query cost.

## Attempt 3 : Index on Temperature and Rating

```
-----+
| -> Sort: AverageRating DESC (actual time=1489.189..1489.190 rows=6 loops=1)
|   -> Filter: ((avg(t.Temperature) > (select #2)) and (avg(r.Rating) > (select #3))) (actual time=1489.148..1489.159 rows=6 loops=1)
|     -> Table scan on <temporary> (actual time=1485.852..1485.866 rows=28 loops=1)
|       -> Aggregate using temporary table (actual time=1485.844..1485.844 rows=28 loops=1)
|         -> Nested loop inner join (cost=15151.34 rows=28624) (actual time=0.576..762.701 rows=229700 loops=1)
|           -> Nested loop inner join (cost=5132.90 rows=5624) (actual time=0.482..50.561 rows=21200 loops=1)
|             -> Nested loop inner join (cost=1399.25 rows=3100) (actual time=0.447..21.912 rows=3100 loops=1)
|               -> Covering index scan on t using idx_temperature (cost=314.25 rows=3100) (actual time=0.405..3.066 rows=310
0 loops=1)
|                 -> Single-row index lookup on c using PRIMARY (Name=t.CityName) (cost=0.25 rows=1) (actual time=0.006..0.006
rows=1 loops=3100)
|               -> Covering index lookup on h using idx_hotel_cityname (CityName=t.CityName) (cost=1.02 rows=2) (actual time=0.0
05..0.008 rows=7 loops=3100)
|             -> Index lookup on r using HotelName (HotelName=h.'Name') (cost=1.27 rows=5) (actual time=0.013..0.033 rows=11 loops
=21200)
|           -> Select #2 (subquery in condition; run only once)
|             -> Aggregate: avg(Temperature.Temperature) (cost=624.25 rows=1) (actual time=0.855..0.855 rows=1 loops=1)
|               -> Covering index scan on Temperature using idx_temperature (cost=314.25 rows=3100) (actual time=0.037..0.660 rows=3100
loops=1)
|           -> Select #3 (subquery in condition; run only once)
|             -> Aggregate: avg(Review.Rating) (cost=1802.75 rows=1) (actual time=2.393..2.393 rows=1 loops=1)
|               -> Covering index scan on Review using idx_rating (cost=953.75 rows=8490) (actual time=0.029..1.858 rows=8492 loops=1)
|
+-----+
```

Likewise, we anticipate enhancing performance by adding an index to attributes present in the HAVING clause, such as the attribute 'Temperature' and 'Rating' in this context. However, we have not observed a reduction in query cost. Similarly, we believe that the size of our table may not be substantial enough to exhibit significant improvements.

## Advanced Query 2

**Advanced features:** join multiple relations, aggregation via GROUP BY  
SELECT

```
    Hotel.Name,  
    Hotel.CityName,  
    Address,  
    ROUND(AVG(Rating), 2) AS Rating  
FROM Hotel  
    JOIN Review ON Hotel.Name = Review.HotelName AND Hotel.CityName =  
Review.CityName WHERE YEAR(Date) > 2015  
GROUP BY  
    Hotel.Name,  
    Hotel.CityName  
HAVING Rating > 3  
ORDER BY Rating DESC  
LIMIT 15;
```

**Top 15 results:**

```
mysql> SELECT Hotel.Name, Hotel.CityName, Address, ROUND(AVG(Rating), 2) AS Rating FROM Hotel JOIN Review ON Hotel.Name = Review.HotelName AND Hotel.CityName = Review.CityName WHERE YEAR(Date) > 2015 GROUP BY Hotel.Name, Hotel.CityName HAVING Rating > 3 ORDER BY Rating DESC LIMIT 15;
```

Name	CityName	Address	Rating
250 Main Hotel	Rockland	250 Main St	5
AC Hotel by Marriott Boston Downtown	Boston	225 Albany Street	5
AC Hotel Miami Beach	Miami Beach	2912 Collins Ave	5
Aloft Bollingbrook	Bollingbrook	500 James Ave	5
Aloft Greenville Downtown	Greenville	5 N Laurens St	5
Aloft Philadelphia Downtown	Philadelphia	101 N Broad St	5
Aloft Sarasota	Sarasota	1401 Ringling Blvd	5
Americas Best Value Inn & Suites-eureka	Eureka	129 4th St	5
Arizona Inn Suites	Yuma	2655 E 4th Ave	5
Hardesono	Yountville	6526 Yount St	5
Basecamp Hotel	South Lake Tahoe	4143 Cedar Ave	5
Baymont Inn Suites - Tullahoma	Tullahoma	2113 N Jackson St	5
Bendel Executive Suites	Lafayette	213 Bendel Rd	5
Best Western Beacon Inn	Grand Haven	1525 E Beacon Blvd	5
Best Western Blackwell Inn	Blackwell	4545 White Ave S	5

15 rows in set (0.04 sec)

**Command:**

```
mysql> Explain Analyze (SELECT  
-> Hotel.Name,  
-> Hotel.CityName,  
-> Address,  
-> ROUND(  
->     AVG(Rating),  
->     2  
-> ) AS Rating  
-> FROM  
-> Hotel  
-> JOIN Review ON Hotel.Name = Review.HotelName  
-> AND Hotel.CityName = Review.CityName  
-> WHERE  
->     YEAR(Date) > 2015  
-> GROUP BY  
->     Hotel.Name,  
->     Hotel.CityName  
-> HAVING  
->     Rating > 3  
-> ORDER BY  
->     Rating DESC  
-> LIMIT  
->     15);
```

## Analysis before indexing:

```
+-----+
|
+-----+
| -> Limit: 15 row(s) (actual time=42.328..42.330 rows=15 loops=1)
|   -> Sort: Rating DESC (actual time=42.327..42.329 rows=15 loops=1)
|     -> Filter: (Rating > 3) (actual time=0.147..41.883 rows=1085 loops=1)
|       -> Stream results (cost=3501.52 rows=93) (actual time=0.145..41.634 rows=1319 loops=1)
|         -> Group aggregate: avg(Review.Rating) (cost=3501.52 rows=93) (actual time=0.142..40.799 rows=1319 loops=1)
|           -> Nested loop inner join (cost=3492.25 rows=93) (actual time=0.127..38.229 rows=4879 loops=1)
|             -> Index scan on Hotel using PRIMARY (cost=189.40 rows=1854) (actual time=0.068..1.114 rows=1854 loops=1)
|               -> Filter: (year(Review.Date) > 2015) (cost=1.27 rows=0.05) (actual time=0.013..0.020 rows=3 loops=1854)
|                 -> Index lookup on Review using HotelName (HotelName=Hotel.Name), with index condition: (Review.CityName = Hotel.CityName) (cost=1.27 rows=5) (actual time=0.012..0.019 rows=5 loops=1854)
|
+-----+
|
+-----+
1 row in set (0.04 sec)
```

## Attempt 1 : Index on Review.Rating

```
+-----+
|
+-----+
| -> Limit: 15 row(s) (actual time=44.512..44.515 rows=15 loops=1)
|   -> Sort: Rating DESC (actual time=44.512..44.513 rows=15 loops=1)
|     -> Filter: (Rating > 3) (actual time=0.134..44.075 rows=1085 loops=1)
|       -> Stream results (cost=3501.52 rows=93) (actual time=0.131..43.805 rows=1319 loops=1)
|         -> Group aggregate: avg(Review.Rating) (cost=3501.52 rows=93) (actual time=0.128..43.004 rows=1319 loops=1)
|           -> Nested loop inner join (cost=3492.25 rows=93) (actual time=0.112..40.369 rows=4879 loops=1)
|             -> Index scan on Hotel using PRIMARY (cost=189.40 rows=1854) (actual time=0.059..0.946 rows=1854 loops=1)
|               -> Filter: (year(Review.Date) > 2015) (cost=1.27 rows=0.05) (actual time=0.013..0.021 rows=3 loops=1854)
|                 -> Index lookup on Review using HotelName (HotelName=Hotel.Name), with index condition: (Review.CityName = Hotel.CityName) (cost=1.27 rows=5) (actual time=0.013..0.020 rows=5 loops=1854)
|
+-----+
|
+-----+
1 row in set (0.05 sec)
```

Likewise, we anticipate enhancing performance by adding an index to attributes present in the HAVING clause, such as the attribute 'Rating' in this context. However, we have not observed a reduction in query cost. Similarly, we believe that the size of our table may not be substantial enough to exhibit significant improvements.

## Attempt 2: Index on Review.Date

```
+-----+
|
+-----+
| -> Limit: 15 row(s) (actual time=41.655..41.658 rows=15 loops=1)
|   -> Sort: Rating DESC (actual time=41.654..41.656 rows=15 loops=1)
|     -> Filter: (Rating > 3) (actual time=0.163..41.197 rows=1085 loops=1)
|       -> Stream results (cost=3501.52 rows=93) (actual time=0.160..40.955 rows=1319 loops=1)
|         -> Group aggregate: avg(Review.Rating) (cost=3501.52 rows=93) (actual time=0.158..40.164 rows=1319 loops=1)
|           -> Nested loop inner join (cost=3492.25 rows=93) (actual time=0.142..37.602 rows=4879 loops=1)
|             -> Index scan on Hotel using PRIMARY (cost=189.40 rows=1854) (actual time=0.075..0.911 rows=1854 loops=1)
|               -> Filter: (year(Review.Date) > 2015) (cost=1.27 rows=0.05) (actual time=0.013..0.019 rows=3 loops=1854)
|                 -> Index lookup on Review using HotelName (HotelName=Hotel.Name), with index condition: (Review.CityName = Hotel.CityName) (cost=1.27 rows=5) (actual time=0.012..0.019 rows=5 loops=1854)
|
+-----+
|
+-----+
1 row in set (0.05 sec)
```

Similarly, we expected to boost performance by indexing attributes mentioned in the WHERE clause, like 'Date'. The logic behind this strategy mirrors that of indexing attributes in the HAVING clause. Despite these efforts, there has been no noticeable reduction in query costs. Once again, it appears that the relatively small size of our table may not allow for the significant improvements we anticipated.

```
+-----+
|
+-----+
| -> Limit: 15 row(s) (actual time=41.620..41.623 rows=15 loops=1)
|   -> Sort: Rating DESC (actual time=41.619..41.621 rows=15 loops=1)
|     -> Filter: (Rating > 3) (actual time=0.123..41.161 rows=1085 loops=1)
|       -> Stream results (cost=3501.52 rows=93) (actual time=0.120..40.931 rows=1319 loops=1)
|         -> Group aggregate: avg(Rewiew.Rating) (cost=3501.52 rows=93) (actual time=0.118..40.145 rows=1319 loops=1)
|           -> Nested loop inner join (cost=3492.25 rows=93) (actual time=0.103..37.722 rows=1879 loops=1)
|             -> Index scan on Hotel using PRIMARY (cost=189.40 rows=1854) (actual time=0.055..0.891 rows=1854 loops=1)
|               -> Filter: (year(Rewiew.Date') > 2015) (cost=1.27 rows=0.05) (actual time=0.013..0.020 rows=3 loops=1854)
|                 -> Index lookup on Review using HotelName (HotelName=Hotel.Name'), with index condition: (Review.CityName = Hotel.CityName) (cost=1.27 rows=5) (actual time=0.012..0.019 rows=5 loops=1854)
|
+-----+
|
+-----+
| 1 row in set (0.04 sec)
```

### Advanced Query 3

SELECT

From Hotel H

### Top 15 results:

name	cityName	Address	LivingWage	CrimeRate
Best Western Plus-prairie Inn	Albany	1100 Price Rd Se	26.484166666666667	1332.7999892234802
Ramada Plaza Albany	Albany	3 Watervliet Ave. Ext.	26.484166666666667	1332.7999892234802
Holiday Inn Express Alpharetta - Roswell	Alpharetta	2950 Mansell Rd	26.484166666666667	399.6000053882599
Wingate By Wyndham Alpharetta	Alpharetta	1005 Kingswood Pl	26.484166666666667	399.6000053882599
Ramada-ankeny	Ankeny	133 Se Delaware Ave	26.271666666666665	2917.2999782562256
Best Western Appleton Inn	Appleton	3033 W College Ave	26.837499999999999	964.8999977111816
Copperleaf Hotel	Appleton	300 W College Ave	26.837499999999999	964.8999977111816
Fairfield Inn Appleton	Appleton	132 N Mall Dr	26.837499999999999	964.8999977111816
Comfort Suites Outlet Center	Asheville	890 Brevard Rd	26.797499999999996	2264.600009918213
Americas Best Value Inn	Auburn	170 Center St	24.356666666666666	1507.1000146865845
Auburn Travelodge Inn & Suites	Auburn	9 16th St Nw	24.356666666666666	1507.1000146865845
Hearthsides Village Cottage Motel	Bethlehem	1267 Main St	24.617499999999996	1244.600020647049
Historic Hotel Bethlehem	Bethlehem	437 Main St	24.617499999999996	1244.600020647049
Holiday Inn Express & Suites Bethlehem	Bethlehem	2201 Cherry Ln	24.617499999999996	1244.600020647049
Residence Inn Allentown Bethlehem/Lehigh Valley Airport	Bethlehem	2180 Motel Dr	24.617499999999996	1244.600020647049

15 rows in set (0.00 sec)



## Command:

```
mysql> EXPLAIN ANALYZE
-> SELECT H.name, H.cityName, H.Address, C.LivingWage, C.CrimeRate
-> From Hotel H JOIN
-> (SELECT * FROM City WHERE LivingWage < (SELECT AVG(LivingWage) FROM City) AND
-> CrimeRate < (SELECT AVG(CrimeRate) FROM City)) C
-> ON H.CityName = C.Name;
+-----+
| EXPLAIN
```

## Analysis before indexing:

```
+-----+
| -> Nested loop inner join (cost=123.81 rows=301) (actual time=1.327..3.185 rows=222 loops=1)
|   -> Filter: ((City.LivingWage < (select #3)) and (City.CrimeRate < (select #4))) (cost=18.36 rows=166) (actual time=1.208..1.853 rows=242 loops=1)
|     -> Table scan on City (cost=18.36 rows=1495) (actual time=0.047..0.497 rows=1495 loops=1)
|       -> Select #3 (subquery in condition; run only once)
|         -> Aggregate: avg(City.LivingWage) (cost=300.75 rows=1) (actual time=0.576..0.576 rows=1 loops=1)
|           -> Table scan on City (cost=151.25 rows=1495) (actual time=0.038..0.442 rows=1495 loops=1)
|       -> Select #4 (subquery in condition; run only once)
|         -> Aggregate: avg(City.CrimeRate) (cost=300.75 rows=1) (actual time=0.540..0.540 rows=1 loops=1)
|           -> Table scan on City (cost=151.25 rows=1495) (actual time=0.031..0.422 rows=1495 loops=1)
|     -> Index lookup on H using idx_hotel_cityname (CityName=City.'Name') (cost=0.45 rows=2) (actual time=0.005..0.005 rows=1 loops=242)
|
+-----+
1 row in set (0.01 sec)
```

## Attempt 1: Index on LivingWage

```
+-----+
| -> Nested loop inner join (cost=215.32 rows=267) (actual time=0.556..2.301 rows=222 loops=1)
|   -> Filter: ((City.LivingWage < (select #3)) and (City.CrimeRate < (select #4))) (cost=121.78 rows=147) (actual time=0.530..1.132 rows=242 loops=1)
|     -> Table scan on City (cost=121.78 rows=1495) (actual time=0.049..0.474 rows=1495 loops=1)
|       -> Select #3 (subquery in condition; run only once)
|         -> Aggregate: avg(City.LivingWage) (cost=300.75 rows=1) (actual time=0.422..0.422 rows=1 loops=1)
|           -> Covering index scan on City using idx_livingwage (cost=151.25 rows=1495) (actual time=0.037..0.326 rows=1495 loops=1)
|       -> Select #4 (subquery in condition; run only once)
|         -> Aggregate: avg(City.CrimeRate) (cost=300.75 rows=1) (actual time=0.473..0.473 rows=1 loops=1)
|           -> Table scan on City (cost=151.25 rows=1495) (actual time=0.032..0.370 rows=1495 loops=1)
|     -> Index lookup on H using idx_hotel_cityname (CityName=City.'Name') (cost=0.45 rows=2) (actual time=0.004..0.005 rows=1 loops=242)
|
+-----+
1 row in set (0.01 sec)
```

Similarly, we expected to improve performance by indexing attributes in the WHERE clause, such as 'LivingWage'. Contrary to our expectations, not only did we fail to see a decrease in query cost, but we also experienced a decline in performance. We suspect that the relatively small size of our table is insufficient to demonstrate notable benefits. Furthermore, the additional overhead associated with implementing the index has negatively impacted overall performance.

```
+-----+
|
+-----+
| -> Nested loop inner join (cost=252.52 rows=256) (actual time=0.578..2.493 rows=222 loops=1)
|   -> Filter: ((City.LivingWage < (select #3)) and (City.CrimeRate < (select #4))) (cost=162.79 rows=141) (actual time=0.504..1.233 rows=242 loops=1)
|     -> Index range scan on City using idx_crimerate over (NULL < CrimeRate < 4225.994471099387) (cost=162.79 rows=424) (actual time=0.014..0.664 rows=424 loops=1)
|     -> Select #3 (subquery in condition; run only once)
|       -> Aggregate: avg(City.LivingWage) (cost=300.75 rows=1) (actual time=0.475..0.475 rows=1 loops=1)
|         -> Table scan on City (cost=151.25 rows=1495) (actual time=0.032..0.371 rows=1495 loops=1)
|       -> Select #4 (subquery in condition; run only once)
|         -> Aggregate: avg(City.CrimeRate) (cost=300.75 rows=1) (actual time=0.396..0.396 rows=1 loops=1)
|           -> Covering index scan on City using idx_crimerate (cost=151.25 rows=1495) (actual time=0.039..0.304 rows=1495 loops=1)
|         -> Index lookup on H using idx_hotel_cityname (CityName=City.'Name') (cost=0.45 rows=2) (actual time=0.004..0.005 rows=1 loops=242)
|
+-----+
|
+-----+
1 row in set (0.01 sec)
```

[illegible]

### Acknowledgements

SELECT

H Name AS HotelName

H.Address.

H.CityName.

C.State.

## CASE

```
WHEN P.Month IN (3, 4, 5) THEN 'Spring'
WHEN P.Month IN (6, 7, 8) THEN 'Summer'
WHEN P.Month IN (9, 10, 11) THEN 'Fall'
ELSE 'Winter'
```

END AS Season,

ROUND(AVG(P.Precipitation), 3) AS AvgPrecipitation

FROM Hotel H

JOIN City C ON H.CityName = C.Name

JOIN Precipitation P ON C.State = P.CityState

GROUP BY

H.Name,  
H.Address,  
H.CityName,  
C.State,  
Season

ORDER BY

H.Name, Season

LIMIT 15;

### Top 15 results:

HotelName	Address	CityName	State	Season	AvgPrecipitation
AC Hotel by Marriott Boston Downtown	225 Albany Street	Boston	Massachusetts	Fall	3.783
AC Hotel by Marriott Boston Downtown	225 Albany Street	Boston	Massachusetts	Spring	3.83
AC Hotel by Marriott Boston Downtown	225 Albany Street	Boston	Massachusetts	Summer	3.453
AC Hotel by Marriott Boston Downtown	225 Albany Street	Boston	Massachusetts	Winter	3.453
AC Hotel Chicago Downtown	630 North Rush Street	Chicago	Illinois	Fall	3.323
AC Hotel Chicago Downtown	630 North Rush Street	Chicago	Illinois	Spring	3.49
AC Hotel Chicago Downtown	630 North Rush Street	Chicago	Illinois	Summer	4.02
AC Hotel Chicago Downtown	630 North Rush Street	Chicago	Illinois	Winter	2.18
AC Hotel Miami Beach	2912 Collins Ave	Miami Beach	Florida	Fall	3.807
AC Hotel Miami Beach	2912 Collins Ave	Miami Beach	Florida	Spring	4.157
AC Hotel Miami Beach	2912 Collins Ave	Miami Beach	Florida	Summer	7.417
AC Hotel Miami Beach	2912 Collins Ave	Miami Beach	Florida	Winter	4.363
Ace Hotel Chicago	311 North Morgan Street	Chicago	Illinois	Fall	3.323
Ace Hotel Chicago	311 North Morgan Street	Chicago	Illinois	Spring	3.49
Ace Hotel Chicago	311 North Morgan Street	Chicago	Illinois	Summer	4.02

### Analysis before indexing:

```
-----+-----  
| -> Limit: 20 row(s)   (actual time=34.205..34.211 rows=20 loops=1)  
|      -> Sort: H.Name , Season, limit input to 20 row(s) per chunk   (actual time=34.204..34.209 rows=20 loops=1)  
|      -> Table scan on <temporary>   (actual time=32.160..32.166 rows=2996 loops=1)  
|      -> Aggregate using temporary table   (actual time=32.166..32.166 rows=2996 loops=1)  
|      -> Nested loop inner join   (cost=3529.66 rows=22210)   (actual time=0.106..0.415 rows=8988 loops=1)  
|          -> Nested loop inner join   (cost=838.30 rows=1854)   (actual time=0.088..4.248 rows=749 loops=1)  
|              -> Table scan on H   (cost=189.40 rows=1854)   (actual time=0.059..0.839 rows=1854 loops=1)  
|                  -> Filter: (C.State is not null)   (cost=0.25 rows=1)   (actual time=0.002..0.002 rows=0 loops=1854)  
|                      -> Single-row index lookup on C using PRIMARY (Name=H.CityName)   (cost=0.25 rows=1)   (actual time=0.001..0.001 rows=1 loops=1854)  
|                          -> Index lookup on P using PRIMARY (CityState=C.State)   (cost=0.25 rows=12)   (actual time=0.005..0.007 rows=12 loops=749)  
|  
+-----+-----
```

## Attempt 1: Index on Precipitation

```
-----+
| -> Limit: 20 row(s) (actual time=33.356..33.362 rows=20 loops=1)
|   -> Sort: H.'Name', Season, limit input to 20 row(s) per chunk (actual time=33.355..33.359 rows=20 loops=1)
|     -> Table scan on <temporary> (actual time=31.381..32.194 rows=2996 loops=1)
|       -> Aggregate using temporary table (actual time=31.378..31.378 rows=2996 loops=1)
|         -> Nested loop inner join (cost=3529.66 rows=22210) (actual time=0.102..10.047 rows=8988 loops=1)
|           -> Nested loop inner join (cost=838.30 rows=1854) (actual time=0.087..3.989 rows=749 loops=1)
|             -> Table scan on H (cost=189.40 rows=1854) (actual time=0.059..0.819 rows=1854 loops=1)
|               -> Filter: (C.State is not null) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=0 loops=1854)
|                 -> Single-row index lookup on C using PRIMARY (Name=H.CityName) (cost=0.25 rows=1) (actual time=0.001..0.001 rows=1 loops=1854)
|                   -> Index lookup on P using PRIMARY (CityState=C.State) (cost=0.25 rows=12) (actual time=0.005..0.007 rows=12 loops=749)
|
|-----+
```

The addition of an index on the attribute present in the SELECT clause, specifically 'Precipitation,' was not expected to influence the cost. The outcomes of our analysis align with these expectations, confirming that the index has no impact on query cost.

## Attempt 2: Index on Address

```
-----+
| -> Limit: 20 row(s) (actual time=62.699..62.706 rows=20 loops=1)
|   -> Sort: H.'Name', Season, limit input to 20 row(s) per chunk (actual time=62.697..62.703 rows=20 loops=1)
|     -> Table scan on <temporary> (actual time=59.049..60.333 rows=2996 loops=1)
|       -> Aggregate using temporary table (actual time=59.045..59.045 rows=2996 loops=1)
|         -> Nested loop inner join (cost=3529.66 rows=22210) (actual time=0.103..19.126 rows=8988 loops=1)
|           -> Nested loop inner join (cost=838.30 rows=1854) (actual time=0.084..7.319 rows=749 loops=1)
|             -> Covering index scan on H using idx_hotel_address (cost=189.40 rows=1854) (actual time=0.061..1.884 rows=1854 loops=1)
|               -> Filter: (C.State is not null) (cost=0.25 rows=1) (actual time=0.003..0.003 rows=0 loops=1854)
|                 -> Single-row index lookup on C using PRIMARY (Name=H.CityName) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=1854)
|                   -> Index lookup on P using PRIMARY (CityState=C.State) (cost=0.25 rows=12) (actual time=0.010..0.014 rows=12 loops=749)
|
|-----+
```

We aimed to boost performance by indexing attributes in the GROUP BY clause, like 'Address'. The idea was that indexing would quickly group identical values, enhancing efficiency. Yet, we haven't seen a decrease in query costs. It seems the scale of our table might not be large enough to manifest significant improvements, given that the advantages of indexing are counterbalanced by the implementation's overhead.

## Attempt 3: Index on 'Precipitation' and 'Address'

```
-----+
| -> Limit: 20 row(s) (actual time=33.245..33.249 rows=20 loops=1)
|   -> Sort: H.'Name', Season, limit input to 20 row(s) per chunk (actual time=33.244..33.247 rows=20 loops=1)
|     -> Table scan on <temporary> (actual time=31.230..32.038 rows=2996 loops=1)
|       -> Aggregate using temporary table (actual time=31.226..31.226 rows=2996 loops=1)
|         -> Nested loop inner join (cost=3529.66 rows=22210) (actual time=0.111..9.810 rows=8988 loops=1)
|           -> Nested loop inner join (cost=838.30 rows=1854) (actual time=0.095..3.868 rows=749 loops=1)
|             -> Covering index scan on H using idx_hotel_address (cost=189.40 rows=1854) (actual time=0.060..0.738 rows=1854 loops=1)
|               -> Filter: (C.State is not null) (cost=0.25 rows=1) (actual time=0.001..0.002 rows=0 loops=1854)
|                 -> Single-row index lookup on C using PRIMARY (Name=H.CityName) (cost=0.25 rows=1) (actual time=0.001..0.001 rows=1 loops=1854)
|                   -> Index lookup on P using PRIMARY (CityState=C.State) (cost=0.25 rows=12) (actual time=0.005..0.007 rows=12 loops=749)
|
|-----+
```

Creating a composite index on 'Precipitation' and 'Address' has yielded the same results as previously observed.